

Prompt Engineering Best Practices

A practical field guide to getting reliable output from Claude, organized around a handful of techniques that consistently matter more than the rest.

1. Introduction

Prompt engineering is closer to writing a clear brief for a new collaborator than it is to programming. The model responds best to instructions that are explicit about the task, the constraints, and the shape of the expected output.

2. Be Clear and Direct

State the task plainly instead of hinting at it. Give context about why the task matters and who the output is for -- it helps the model calibrate tone and level of detail. Specify the output format explicitly rather than leaving it implicit; asking for a numbered list gets a numbered list, asking for an answer gets prose that may or may not be structured the way you wanted.

3. Use Examples

Multishot prompting -- including a small number of input and output examples in the prompt -- is one of the most reliable ways to improve output quality and consistency, especially for tasks with a specific format or edge cases that are hard to describe in the abstract. Three to five diverse examples usually outperform a long written description of the rules.

4. Encourage Step by Step Reasoning

For tasks with multiple steps or that require weighing tradeoffs, asking the model to think through the problem before giving a final answer improves accuracy, particularly on tasks involving arithmetic, multi-step logic, or synthesizing several pieces of context.

5. Common Pitfalls

Vague instructions produce vague output -- if a prompt would leave a new team member unsure what to do, it will leave the model unsure too. Overloading a single prompt with many unrelated tasks tends to produce worse results on all of them than splitting the work into focused steps. And omitting the output format spec is the single most common cause of an otherwise-correct answer arriving in the wrong shape.